

# CHAPTER 7

## More SQL: Complex Queries, Triggers, Views, and Schema Modification

# Chapter 7 Outline

- ▶ More Complex SQL Retrieval Queries
- ▶ Specifying Semantic Constraints as Assertions and Actions as Triggers
- ▶ Views (Virtual Tables) in SQL
- ▶ Schema Modification in SQL

# More Complex SQL Retrieval Queries

- ▶ Additional features allow users to specify more complex retrievals from database:
  - ▶ Nested queries, joined tables, and outer joins (in the FROM clause), aggregate functions, and grouping

# Comparisons Involving NULL and Three-Valued Logic

- ▶ Meanings of `NULL`
  - ▶ Unknown value
  - ▶ Unavailable or withheld value
  - ▶ Not applicable attribute
- ▶ Each individual `NULL` value considered to be different from every other `NULL` value
- ▶ SQL uses a three-valued logic:
  - ▶ `TRUE`, `FALSE`, and `UNKNOWN` (like *Maybe*)
- ▶ `NULL = NULL` comparison is avoided

# Comparisons Involving NULL and Three-Valued Logic (cont'd.)

**Table 7.1** Logical Connectives in Three-Valued Logic

(a)	<b>AND</b>	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	FALSE	UNKNOWN
	FALSE	FALSE	FALSE	FALSE
	UNKNOWN	UNKNOWN	FALSE	UNKNOWN
(b)	<b>OR</b>	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c)	<b>NOT</b>			
	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		

# Comparisons Involving NULL and Three-Valued Logic (cont'd.)

- ▶ SQL allows queries that check whether an attribute value is `NULL`
  - ▶ `IS` or `IS NOT NULL`

**Query 18.** Retrieve the names of all employees who do not have supervisors.

```
Q18:  SELECT  Fname, Lname
      FROM    EMPLOYEE
      WHERE   Super_ssn IS NULL;
```

# Nested Queries, Tuples, and Set/Multiset Comparisons

- ▶ **Nested queries**
  - ▶ Complete select-from-where blocks within WHERE clause of another query
  - ▶ Outer query and nested subqueries
- ▶ **Comparison operator `IN`**
  - ▶ Compares value  $v$  with a set (or multiset) of values  $V$
  - ▶ Evaluates to `TRUE` if  $v$  is one of the elements in  $V$

# Nested Queries (cont'd.)

```
Q4A:  SELECT DISTINCT Pnumber
      FROM PROJECT
      WHERE Pnumber IN
        ( SELECT Pnumber
          FROM PROJECT, DEPARTMENT, EMPLOYEE
          WHERE Dnum=Dnumber AND
                Mgr_ssn=Ssn AND Lname='Smith' )

      OR
      Pnumber IN
        ( SELECT Pno
          FROM WORKS_ON, EMPLOYEE
          WHERE Essn=Ssn AND Lname='Smith' );
```



# Nested Queries (cont'd.)

- ▶ Use tuples of values in comparisons
  - ▶ Place them within parentheses

```
SELECT    DISTINCT Essn
FROM      WORKS_ON
WHERE     (Pno, Hours) IN ( SELECT    Pno, Hours
                           FROM      WORKS_ON
                           WHERE     Essn='123456789' );
```

# Nested Queries (cont'd.)

- ▶ Use other comparison operators to compare a single value  $v$ 
  - ▶  $=$  ANY (or  $=$  SOME) operator
    - ▶ Returns TRUE if the value  $v$  is equal to some value in the set  $V$  and is hence equivalent to IN
  - ▶ Other operators that can be combined with ANY (or SOME):  $>$ ,  $>=$ ,  $<$ ,  $<=$ , and  $<>$
  - ▶ ALL: value must exceed all values from nested query

```
SELECT  Lname, Fname
FROM    EMPLOYEE
WHERE   Salary > ALL ( SELECT  Salary
                        FROM    EMPLOYEE
                        WHERE   Dno=5 );
```

# Nested Queries (cont'd.)

- ▶ Avoid potential errors and ambiguities
  - ▶ Create tuple variables (aliases) for all tables referenced in SQL query

**Query 16.** Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
Q16:  SELECT    E.Fname, E.Lname
      FROM      EMPLOYEE AS E
      WHERE     E.Ssn IN ( SELECT    Essn
                          FROM      DEPENDENT AS D
                          WHERE     E.Fname=D.Dependent_name
                          AND E.Sex=D.Sex );
```

# Correlated Nested Queries

- ▶ Queries that are nested using the = or IN comparison operator can be collapsed into one single block: E.g., Q16 can be written as:
- ▶ **Q16A:   SELECT       E.Fname, E.Lname  
          FROM       EMPLOYEE AS E, DEPENDENT AS D  
          WHERE       E.Ssn=D.Essn AND E.Sex=D.Sex  
                      AND  
          E.Fname=D.Dependent\_name;**
- ▶ **Correlated** nested query
  - ▶ Evaluated once for each tuple in the outer query

# The EXISTS and UNIQUE Functions in SQL for correlating queries

- ▶ `EXISTS` function
  - ▶ Check whether the result of a correlated nested query is empty or not. They are Boolean functions that return a `TRUE` or `FALSE` result.
- ▶ `EXISTS` and `NOT EXISTS`
  - ▶ Typically used in conjunction with a correlated nested query
- ▶ SQL function `UNIQUE (Q)`
  - ▶ Returns `TRUE` if there are no duplicate tuples in the result of query `Q`

# USE of EXISTS

**Q7:**

```
SELECT Fname, Lname
FROM Employee
WHERE EXISTS (SELECT *
               FROM DEPENDENT
               WHERE Ssn= Essn)

      AND EXISTS (SELECT *
                  FROM Department
                  WHERE Ssn= Mgr_Ssn)
```

# USE OF NOT EXISTS

To achieve the “for all” (universal quantifier) effect, we use double negation this way in SQL:

Query: List first and last name of employees who work on ALL projects controlled by Dno=5.

```
SELECT Fname, Lname
FROM Employee
WHERE NOT EXISTS ( (SELECT Pnumber
                     FROM PROJECT
                     WHERE Dno=5)

                   EXCEPT (SELECT Pno
                              FROM WORKS_ON
                              WHERE Ssn= ESsn)
```

The above is equivalent to double negation: List names of those employees for whom there does NOT exist a project managed by department no. 5 that they do NOT work on.

# Double Negation to accomplish “for all” in SQL

```
► Q3B:  SELECT      Lname, Fname
        FROM        EMPLOYEE
        WHERE NOT EXISTS (      SELECT      *
                                FROM WORKS_ON B
                                WHERE      ( B.Pno IN ( SELECT Pnumber
                                                         FROM PROJECT
                                                         WHERE Dnum=5
                                                         AND
                                                         NOT EXISTS (SELECT *
                                                                    FROM WORKS_ON C
                                                                    WHERE C.Essn=Ssn
                                                                    AND      C.Pno=B.Pno )));
```

The above is a direct rendering of: List names of those employees for whom there does NOT exist a project managed by department no. 5 that they do NOT work on.



# Explicit Sets and Renaming of Attributes in SQL

- ▶ Can use explicit set of values in WHERE clause

```
Q17:    SELECT DISTINCT Essn
        FROM      WORKS_ON
        WHERE Pno IN (1, 2, 3);
```

- ▶ Use qualifier AS followed by desired new name
  - ▶ Rename any attribute that appears in the result of a query

```
Q8A:    SELECT      E.Lname AS Employee_name, S.Lname AS Supervisor_name
        FROM      EMPLOYEE AS E, EMPLOYEE AS S
        WHERE      E.Super_ssn=S.Ssn;
```

# Specifying Joined Tables in the FROM Clause of SQL

- ▶ **Joined table**

- ▶ Permits users to specify a table resulting from a join operation in the FROM clause of a query

- ▶ **The FROM clause in Q1A**

- ▶ Contains a single joined table. JOIN may also be called INNER JOIN

```
Q1A:  SELECT  Fname, Lname, Address
      FROM    (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
      WHERE   Dname='Research';
```

# Different Types of JOINed Tables in SQL

- ▶ Specify different types of join
  - ▶ NATURAL JOIN
  - ▶ Various types of OUTER JOIN (LEFT, RIGHT, FULL )
- ▶ NATURAL JOIN on two relations R and S
  - ▶ No join condition specified
  - ▶ Is equivalent to an implicit EQUIJOIN condition for each pair of attributes with same name from R and S

# NATURAL JOIN

- Rename attributes of one relation so it can be joined with another using NATURAL JOIN:

```
Q1B:      SELECT      Fname, Lname, Address
           FROM (EMPLOYEE NATURAL JOIN
                 (DEPARTMENT AS DEPT (Dname, Dno, Mssn,
                                       Msdate)))
           WHERE      Dname='Research';
```

The above works with  $EMPLOYEE.Dno = DEPT.Dno$  as an implicit join condition

# INNER and OUTER Joins

- ▶ INNER JOIN (**versus** OUTER JOIN)
  - ▶ Default type of join in a joined table
  - ▶ Tuple is included in the result only if a matching tuple exists in the other relation
- ▶ LEFT OUTER JOIN
  - ▶ Every tuple in left table must appear in result
  - ▶ If no matching tuple
    - ▶ Padded with NULL values for attributes of right table
- ▶ RIGHT OUTER JOIN
  - ▶ Every tuple in right table must appear in result
  - ▶ If no matching tuple
    - ▶ Padded with NULL values for attributes of left table

# Example: LEFT OUTER JOIN

```
SELECT E.Lname AS Employee_Name  
       S.Lname AS Supervisor_Name
```

```
FROM Employee AS E LEFT OUTER JOIN EMPLOYEE AS S  
      ON E.Super_ssn = S.Ssn)
```

## ALTERNATE SYNTAX:

```
SELECT E.Lname , S.Lname  
FROM EMPLOYEE E, EMPLOYEE S  
WHERE E.Super_ssn + = S.Ssn
```

# Multiway JOIN in the FROM clause

- ▶ FULL OUTER JOIN - combines result if LEFT and RIGHT OUTER JOIN
- ▶ Can nest JOIN specifications for a multiway join:

**Q2A:**     **SELECT** Pnumber, Dnum, Lname, Address,  
          Bdate

**FROM**     ((PROJECT JOIN DEPARTMENT ON  
              Dnum=Dnumber) JOIN EMPLOYEE ON  
              Mgr\_ssn=Ssn)

**WHERE**   Plocation='Stafford';

# Aggregate Functions in SQL

- ▶ Used to summarize information from multiple tuples into a single-tuple summary
- ▶ Built-in aggregate functions
  - ▶ `COUNT`, `SUM`, `MAX`, `MIN`, and `AVG`
- ▶ **Grouping**
  - ▶ Create subgroups of tuples before summarizing
- ▶ To select entire groups, `HAVING` clause is used
- ▶ Aggregate functions can be used in the `SELECT` clause or in a `HAVING` clause



# Renaming Results of Aggregation

- ▶ Following query returns a single row of computed values from EMPLOYEE table:

```
Q19:      SELECT  SUM (Salary), MAX (Salary), MIN  
           (Salary), AVG           (Salary)  
           FROM    EMPLOYEE;
```

- ▶ The result can be presented with new names:

```
Q19A:     SELECT  SUM (Salary) AS Total_Sal, MAX (Salary)  
           AS      Highest_Sal, MIN (Salary) AS Lowest_Sal,  
           AVG     (Salary) AS Average_Sal  
           FROM    EMPLOYEE;
```

# Aggregate Functions in SQL (cont'd.)

- ▶ NULL values are discarded when aggregate functions are applied to a particular column

**Query 20.** Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
Q20:  SELECT    SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)
      FROM      (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
      WHERE     Dname='Research';
```

**Queries 21 and 22.** Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).

```
Q21:  SELECT    COUNT (*)
      FROM      EMPLOYEE;
```

```
Q22:  SELECT    COUNT (*)
      FROM      EMPLOYEE, DEPARTMENT
      WHERE     DNO=DNUMBER AND DNAME='Research';
```

# Aggregate Functions on Booleans

- ▶ SOME and ALL may be applied as functions on Boolean Values.
- ▶ SOME returns true if at least one element in the collection is TRUE (similar to OR)
- ▶ ALL returns true if all of the elements in the collection are TRUE (similar to AND)

# Grouping: The GROUP BY Clause

- ▶ **Partition** relation into subsets of tuples
  - ▶ Based on **grouping attribute(s)**
  - ▶ Apply function to each such group independently
- ▶ **GROUP BY** clause
  - ▶ Specifies grouping attributes
- ▶ **COUNT (\*)** counts the number of rows in the group

# Examples of GROUP BY

- ▶ The grouping attribute must appear in the SELECT clause:

```
Q24:      SELECT    Dno, COUNT (*), AVG (Salary)
           FROM      EMPLOYEE
           GROUP BY  Dno;
```

- ▶ If the grouping attribute has NULL as a possible value, then a separate group is created for the null value (e.g., null Dno in the above query)
- ▶ GROUP BY may be applied to the result of a JOIN:

```
Q25:  SELECT          Pnumber, Pname, COUNT (*)
      FROM    PROJECT, WORKS_ON
      WHERE    Pnumber=Pno
      GROUP BY  Pnumber, Pname;
```

# Grouping: The GROUP BY and HAVING Clauses (cont'd.)

- ▶ **HAVING** clause
  - ▶ Provides a condition to select or reject an entire group:
- ▶ **Query 26.** For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

```
Q26:    SELECT Pnumber, Pname, COUNT (*)
FROM      PROJECT, WORKS_ON
WHERE Pnumber=Pno
GROUP BY  Pnumber, Pname
HAVING COUNT (*) > 2;
```

# Combining the WHERE and the HAVING Clause

- ▶ Consider the query: we want to count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work.

- ▶ **INCORRECT QUERY:**

```
SELECT      Dno, COUNT (*)  
FROM        EMPLOYEE  
WHERE       Salary>40000  
GROUP BY Dno  
HAVING      COUNT (*) > 5;
```

# Combining the WHERE and the HAVING Clause (continued)

## Correct Specification of the Query:

- Note: the WHERE clause applies tuple by tuple whereas HAVING applies to entire group of tuples

**Query 28.** For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

```
Q28:  SELECT  Dnumber, COUNT (*)
      FROM    DEPARTMENT, EMPLOYEE
      WHERE   Dnumber=Dno AND Salary>40000 AND
             ( SELECT  Dno
               FROM    EMPLOYEE
               GROUP BY Dno
               HAVING   COUNT (*) > 5)
```



# Use of WITH

- ▶ The WITH clause allows a user to define a table that will only be used in a particular query (not available in all SQL implementations)
- ▶ Used for convenience to create a temporary “View” and use that immediately in a query
- ▶ Allows a more straightforward way of looking a step-by-step query

# Example of WITH

- ▶ See an alternate approach to doing Q28:

- ▶ **Q28':**                **WITHBIGDEPTS (Dno) AS**  
                          **( SELECT    Dno**  
                          **FROM      EMPLOYEE**  
                          **GROUP BY Dno**  
                          **HAVING    COUNT (\*) > 5)**  
**SELECT                Dno, COUNT (\*)**  
**FROM                  EMPLOYEE**  
**WHERE                 Salary>40000 AND Dno IN BIGDEPTS**  
**GROUP BY Dno;**

# Use of CASE

- ▶ SQL also has a CASE construct
- ▶ Used when a value can be different based on certain conditions.
- ▶ Can be used in any part of an SQL query where a value is expected
- ▶ Applicable when querying, inserting or updating tuples

# EXAMPLE of use of CASE

- ▶ The following example shows that employees are receiving different raises in different departments (A variation of the update U6)

- ▶ **U6':**  
**UPDATE**      **EMPLOYEE**  
  
**SET**      Salary =  
  
**CASE**      **WHEN** Dno = 5 **THEN**      Salary +  
2000  
  
                 **WHEN** Dno = 4 **THEN**      Salary + 1500  
                 **WHEN** Dno = 1 **THEN**      Salary + 3000

# Recursive Queries in SQL

- ▶ An example of a **recursive relationship** between tuples of the same type is the relationship between an employee and a supervisor.
- ▶ This relationship is described by the foreign key `Super_ssn` of the `EMPLOYEE` relation
- ▶ An example of a **recursive operation** is to retrieve all supervisees of a supervisory employee  $e$  at all levels—that is, all employees  $e'$  directly supervised by  $e$ , all employees  $e''$  directly supervised by each employee  $e'$ , all employees  $e'''$  directly supervised by each employee  $e''$ , and so on. Thus the CEO would have each employee in the company as a supervisee in the resulting table. Example shows such table `SUP_EMP` with 2 columns (`Supervisor, Supervisee(any level)`):

# An EXAMPLE of RECURSIVE Query

► Q29: WITH RECURSIVE SUP\_EMP (SupSsn, EmpSsn) AS

```
SELECT      SupervisorSsn, Ssn
```

```
FROM EMPLOYEE
```

```
UNION
```

```
SELECT      E.Ssn, S.SupSsn
```

```
FROM EMPLOYEE AS E, SUP_EMP AS S
```

```
WHERE      E.SupervisorSsn = S.EmpSsn)
```

```
SELECT      *
```

```
FROM SUP_EMP;
```

- The above query starts with an empty SUP\_EMP and successively builds SUP\_EMP table by computing immediate supervisees first, then second level supervisees, etc. until a **fixed point** is reached and no more supervisees can be added

# EXPANDED Block Structure of SQL Queries

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```

# Specifying Constraints as Assertions and Actions as Triggers

- ▶ Semantic Constraints: The following are beyond the scope of the EER and relational model
- ▶ **CREATE ASSERTION**
  - ▶ Specify additional types of constraints outside scope of built-in relational model constraints
- ▶ **CREATE TRIGGER**
  - ▶ Specify automatic actions that database system will perform when certain events and conditions occur



# Specifying General Constraints as Assertions in SQL

## ► CREATE ASSERTION

- Specify a query that selects any tuples that violate the desired condition
- Use only in cases where it goes beyond a simple CHECK which applies to individual attributes and domains

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK ( NOT EXISTS ( SELECT *
                     FROM   EMPLOYEE E, EMPLOYEE M,
                     WHERE  E.Salary>M.Salary
                           AND E.Dno=D.Dnumber
                           AND D.Mgr_ssn=M.Ssn ) );
```

# Introduction to Triggers in SQL

- ▶ `CREATE TRIGGER` **statement**
  - ▶ Used to monitor the database
- ▶ Typical trigger has three components which make it a rule for an “active database “ (more on active databases) :
  - ▶ **Event(s)**
  - ▶ **Condition**
  - ▶ **Action**

# USE OF TRIGGERS

- ▶ AN EXAMPLE with standard Syntax. (Note : other SQL implementations like PostgreSQL use a different syntax.)

**R5:**

```
CREATE TRIGGER SALARY_VIOLATION  
BEFORE INSERT OR UPDATE OF Salary, Supervisor_ssn ON EMPLOYEE  
  
FOR EACH ROW  
WHEN (NEW.SALARY > ( SELECT Salary FROM EMPLOYEE  
                        WHERE Ssn = NEW. Supervisor_Ssn))  
INFORM_SUPERVISOR (NEW.Supervisor.Ssn, New.Ssn)
```

**Sample Example:** <http://www.mysqltutorial.org/create-the-first-trigger-in-mysql.aspx>