

Views (Virtual Tables) in SQL

- ▶ Concept of a view in SQL
 - ▶ Single table derived from other tables called the **defining tables**
 - ▶ Considered to be a virtual table that is not necessarily populated

Specification of Views in SQL

- ▶ **CREATE VIEW** command

- ▶ Give table name, list of attribute names, and a query to specify the contents of the view

```
V1:  CREATE VIEW  WORKS_ON1
      AS SELECT   Fname, Lname, Pname, Hours
          FROM     EMPLOYEE, PROJECT, WORKS_ON
          WHERE    Ssn=Essn AND Pno=Pnumber;
```

```
V2:  CREATE VIEW  DEPT_INFO(Dept_name, No_of_emps, Total_sal)
      AS SELECT   Dname, COUNT (*), SUM (Salary)
          FROM     DEPARTMENT, EMPLOYEE
          WHERE    Dnumber=Dno
          GROUP BY Dname;
```

Specification of Views in SQL (cont'd.)

- ▶ Once a View is defined, SQL queries can use the View relation in the FROM clause
- ▶ View is always up-to-date
 - ▶ Responsibility of the DBMS and not the user
- ▶ **DROP VIEW** command
 - ▶ Dispose of a view

View Implementation, View Update, and Inline Views

- ▶ Complex problem of efficiently implementing a view for querying
- ▶ **Strategy1: Query modification** approach
 - ▶ Compute the view as and when needed. Do not store permanently
 - ▶ Modify view query into a query on underlying base tables
 - ▶ Disadvantage: inefficient for views defined via complex queries that are time-consuming to execute

View Materialization

- ▶ **Strategy 2: View materialization**
 - ▶ Physically create a temporary view table when the view is first queried
 - ▶ Keep that table on the assumption that other queries on the view will follow
 - ▶ Requires efficient strategy for automatically updating the view table when the base tables are updated
- ▶ **Incremental update strategy for materialized views**
 - ▶ DBMS determines what new tuples must be inserted, deleted, or modified in a materialized view table
 - ▶ Constraints on view updates:
<http://www.informit.com/articles/article.aspx?p=130855&seqNum=4>

View Materialization (contd.)

- ▶ Multiple ways to handle materialization:
 - ▶ **immediate update** strategy updates a view as soon as the base tables are changed
 - ▶ **lazy update** strategy updates the view when needed by a view query
 - ▶ **periodic update** strategy updates the view periodically (in the latter strategy, a view query may get a result that is not up-to-date). This is commonly used in Banks, Retail store operations, etc.

View Update

- ▶ Update on a view defined on a single table without any aggregate functions
 - ▶ Can be mapped to an update on underlying base table- possible if the primary key is preserved in the view
- ▶ Update not permitted on aggregate views. E.g.,

```
UV2: UPDATE      DEPT_INFO
      SET         Total_sal=100000
      WHERE       Dname='Research';
```

cannot be processed because Total_sal is a computed value in the view definition

View Update and Inline Views

- ▶ View involving joins
 - ▶ Often not possible for DBMS to determine which of the updates is intended
- ▶ Clause **WITH CHECK OPTION**
 - ▶ Must be added at the end of the view definition if a view is to be updated to make sure that tuples being updated stay in the view
- ▶ **In-line view**
 - ▶ Defined in the `FROM` clause of an SQL query (e.g., we saw its used in the `WITH` example)

Views as authorization mechanism

- ▶ SQL query authorization statements (GRANT and REVOKE)
- ▶ Views can be used to hide certain attributes or tuples from unauthorized users
- ▶ E.g., For a user who is only allowed to see employee information for those who work for department 5, he may only access the view **DEPT5EMP**:

```
CREATE VIEW    DEPT5EMP  AS
SELECT        *
FROM          EMPLOYEE
WHERE         Dno = 5;
```

Schema Change Statements in SQL

- ▶ **Schema evolution commands**

- ▶ DBA may want to change the schema while the database is operational
- ▶ Does not require recompilation of the database schema

The DROP Command

- ▶ **DROP command**
 - ▶ Used to drop named schema elements, such as tables, domains, or constraint
- ▶ **Drop behavior options:**
 - ▶ `CASCADE` and `RESTRICT`
- ▶ **Example:**
 - ▶ `DROP SCHEMA COMPANY CASCADE;`
 - ▶ This removes the schema and all its elements including tables, views, constraints, etc.

The ALTER table command

- ▶ **Alter table actions include:**

- ▶ Adding or dropping a column (attribute)
- ▶ Changing a column definition
- ▶ Adding or dropping table constraints

- ▶ **Example:**

- ▶ `ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job VARCHAR(12);`

Adding and Dropping Constraints

- ▶ Change constraints specified on a table
 - ▶ Add or drop a named constraint

```
ALTER TABLE COMPANY.EMPLOYEE  
DROP CONSTRAINT EMPSUPERFK CASCADE;
```

Dropping Columns, Default Values

- ▶ To drop a column
 - ▶ Choose either `CASCADE` or `RESTRICT`
 - ▶ `CASCADE` would drop the column from views etc.
`RESTRICT` is possible if no views refer to it.

**ALTER TABLE COMPANY.EMPLOYEE DROP COLUMN
Address CASCADE;**

- ▶ Default values can be dropped and altered :

**ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN
Mgr_ssn DROP DEFAULT;**

**ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN
Mgr_ssn SET DEFAULT '333445555';**

Table 7.2 Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

```
CREATE TABLE <table name> ( <column name> <column type> [ <attribute constraint> ]  
                             { , <column name> <column type> [ <attribute constraint> ] }  
                             [ <table constraint> { , <table constraint> } ] )
```

```
DROP TABLE <table name>  
ALTER TABLE <table name> ADD <column name> <column type>
```

```
SELECT [ DISTINCT ] <attribute list>  
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }  
[ WHERE <condition> ]  
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]  
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]
```

```
<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) )  
                    { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) } )
```

```
<grouping attributes> ::= <column name> { , <column name> }
```

```
<order> ::= ( ASC | DESC )
```

```
INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]  
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) }  
| <select statement> )
```

continued on next slide

Table 7.2 (continued)

Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

DELETE FROM <table name>

[WHERE <selection condition>]

UPDATE <table name>

SET <column name> = <value expression> { , <column name> = <value expression> }

[WHERE <selection condition>]

CREATE [UNIQUE] INDEX <index name>

ON <table name> (<column name> [<order>] { , <column name> [<order>] })

[CLUSTER]

DROP INDEX <index name>

CREATE VIEW <view name> [(<column name> { , <column name> })]

AS <select statement>

DROP VIEW <view name>

NOTE: The commands for creating and dropping indexes are not part of standard SQL.

Summary

- ▶ Complex SQL:
 - ▶ Nested queries, joined tables (in the FROM clause), outer joins, aggregate functions, grouping
- ▶ Handling semantic constraints with `CREATE ASSERTION` and `CREATE TRIGGER`
- ▶ `CREATE VIEW` statement and materialization strategies
- ▶ Schema Modification for the DBAs using `ALTER TABLE`, `ADD` and `DROP COLUMN`, `ALTER CONSTRAINT` etc.