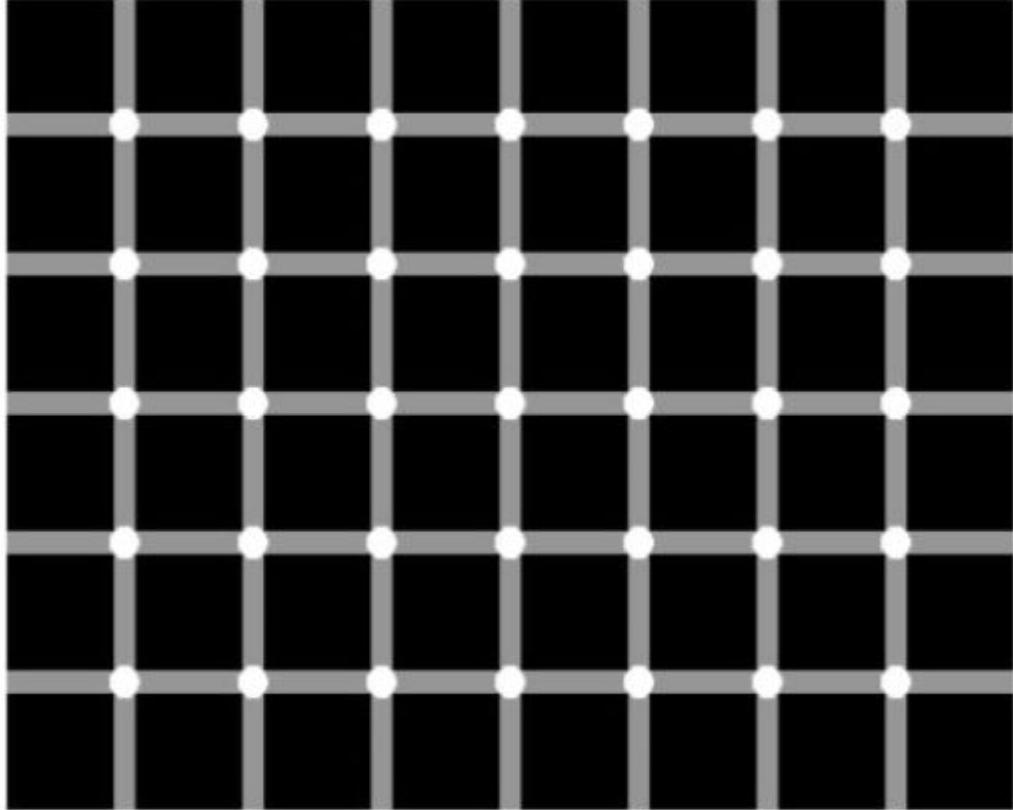




# NoSQL Principles and Industrial Applications

Siddharth Srivastava



# Agenda

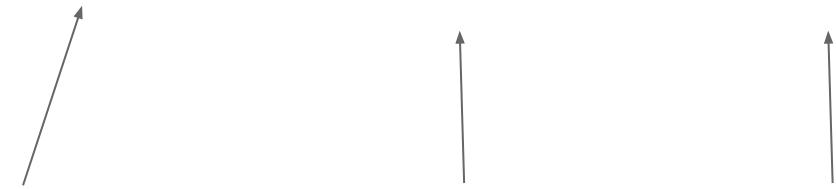
- Current **trends** in data management & computing
- Relational vs. NoSQL databases
  - the value of **relational databases**
  - new **requirements**
  - NoSQL features, strengths and challenges
- **Types** of NoSQL databases
  - **key-value stores**, **document databases**, **column-family databases**, **graph databases**
  - principles and examples

# RDBMS vs NoSQL - Intuition

Query: Who is the faculty for DBMS at IIIT SriCity ?

# RDBMS vs NoSQL: Intuition

Query: Who is the **faculty** for **DBMS** at **IIIT SriCity**



Details of Faculty

Subject

Location

# RDBMS vs NoSQL: Intuition

## RDBMS Approach

- Look for the subject in Subject Table - DBMS
  - (Subject code, Subject Name)
- Get the subject code - DB001
- Get the location code from Location Table - L001
  - (Location Code, Location Name)
- Obtain the faculty code - F001
  - (Subject Code, Location Code, Faculty Code)
- Get Faculty Details - Dr. Prerana Mukherjee

# RDBMS vs NoSQL: Intuition

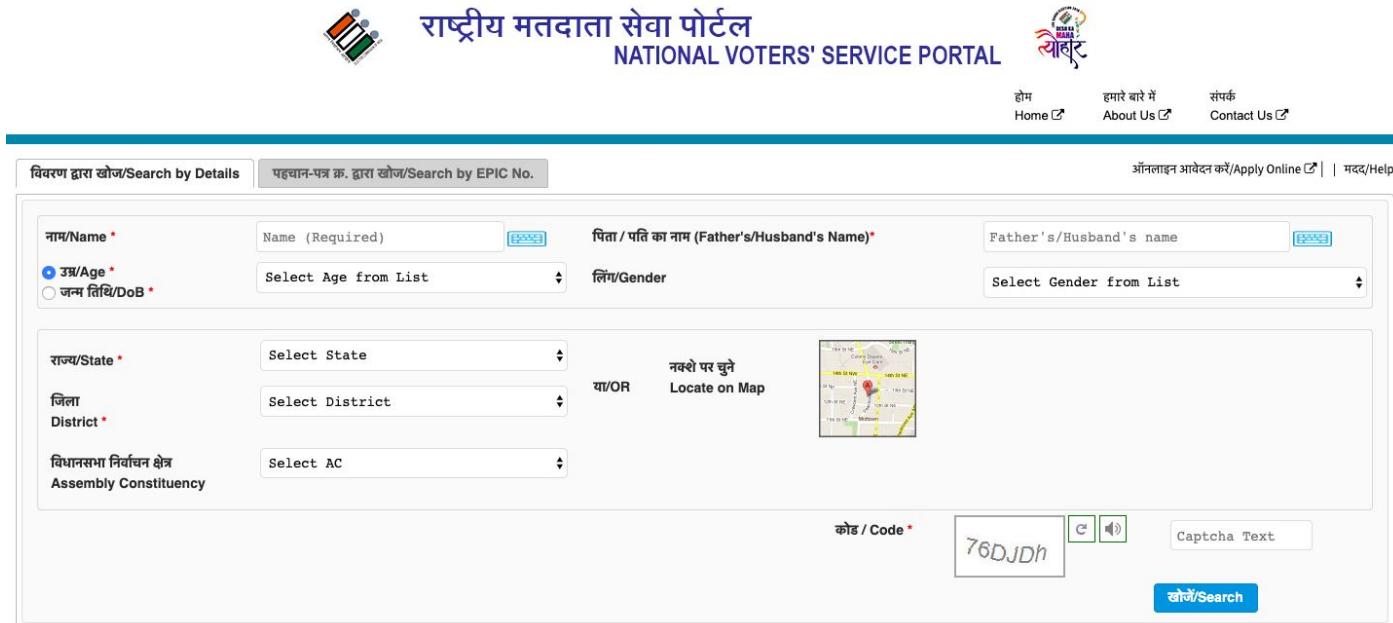
## NoSQL Approach

- Look for **best match** from the **available content** where
  - Subject is 'Database'
  - Location is 'IIIT SriCity'
- Output: Dr. Prerana Mukherjee

(Imagine you are doing CTRL + F in a file)

# RDBMS vs NoSQL: Intuition

Now let us scale up the problem - India's National Electoral Search



The screenshot shows the search interface for the National Voters' Service Portal. It features two main search options: "विवरण द्वारा खोजा/Search by Details" and "पहचान-पत्र क्र. द्वारा खोजा/Search by EPIC No.". Below these are fields for Name, Father's/Husband's Name, Age, Gender, State, District, and Assembly Constituency. A map allows users to "Locate on Map". At the bottom, there is a CAPTCHA field with the code "76DJDh" and a "Search" button.

राष्ट्रीय मतदाता सेवा पोर्टल  
NATIONAL VOTERS' SERVICE PORTAL

विवरण द्वारा खोजा/Search by Details      पहचान-पत्र क्र. द्वारा खोजा/Search by EPIC No.

ऑनलाइन आवेदन करें/Apply Online | मदद/Help

नाम/Name \*  
 उम्र/Age \*  
 जन्म तिथि/DoB \*

Name (Required)  
Select Age from List  
Father's/Husband's name  
Select Gender from List

पिता / पति का नाम (Father's/Husband's Name)\*  
तिक्का/Gender

राज्य/State \*  
ज़िला District \*  
विधानसभा निर्वाचन क्षेत्र Assembly Constituency

Select State  
Select District  
Select AC

या/OR  
नक्शे पर चुने  
Locate on Map

कोड / Code \*  
76DJDh  
Captcha Text

खोजें/Search

# RDBMS vs NoSQL: Intuition

## Challenges (National Electoral Search)

- 80 billion voters
- Each voter has nearly 20 attributes (multilingual name, age, dob, address, polling station, its coordinates etc.)
  - The names across the country are in English and regional languages
  - Lots of missing data
  - Age changes over time
  - Search support with multilingual inputs
  - Location based search - a nightmare !!
- One can update details as well

# RDBMS vs NoSQL: Intuition

# Challenges (National Electoral Search)

- Response time should be less than 50ms
  - Number of concurrent users - about 50 lakhs
  - Data size: 6 TeraBytes

# RDBMS vs NoSQL: Intuition

## How to construct with RDBMS (National Electoral Search)

- Create tables for voter details, constituency details
- Create indexes on all possible search combinations
  - size: 500 GB for one combination, in total 15 search combinations
  - RDBMS have trouble updating large indexes in real time (recall user can update their details)

# RDBMS vs NoSQL: Intuition

## How to construct with RDBMS (National Electoral Search)

- For response time, ensure that the information is available on multiple servers
  - RDBMS are, *in general*, not good at distributing data
- Location based search - name and his location has input, search within a radius
  - search is extremely slow

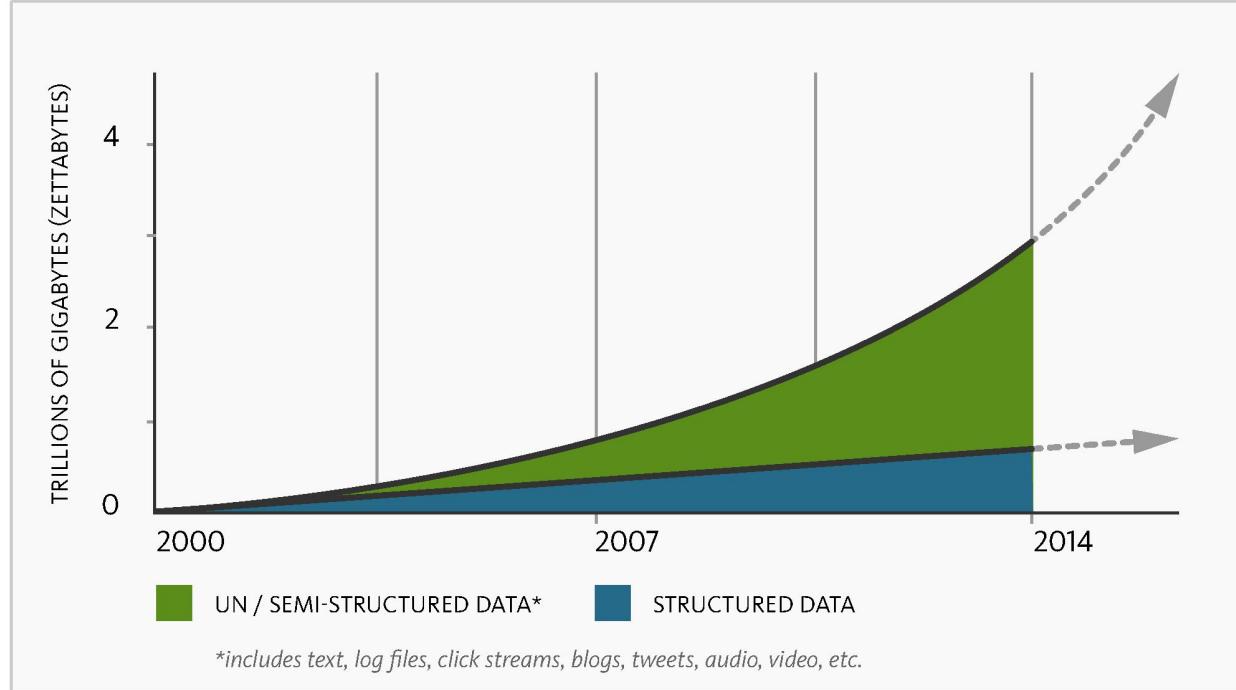
# RDBMS vs NoSQL: Intuition

## What are we looking for (National Electoral Search)

- A solution than can
  - Manage large indexes and update
  - Distribute and search the data on multiple servers
  - Has very low latency
  - Is reliable
  - Supports heterogeneous fields i.e. data is unstructured

NoSQL has all the above properties and many more...

# Current Trends: Big Data



- **Volume, Velocity and Variety of data**

# RDBMS for Big Data

- relational **schema**
  - data in tuples
  - **a priori** known schema
- schema **normalization**
  - data split into tables (3NF)
  - queries merge the data
- **transaction** support
  - trans. management with ACID
  - Atomicity, Consistency, Isolation, Durability
  - safety first
- but current data are naturally **flexible**
- **inefficient** for large data
- slow in **distributed** environment
- **full transactions** very inefficient in **distributed** envir.

# NoSQL Databases

- **What is “NoSQL”?**

- term used in late 90s for a different type of technology:  
Carlo Strozzi: [http://www.strozzi.it/cgi-bin/CSA/tw7/l/en\\_US/NoSQL/](http://www.strozzi.it/cgi-bin/CSA/tw7/l/en_US/NoSQL/)
- “Not Only SQL”?
  - but many RDBMS are also “not just SQL”

“NoSQL is an accidental term with no precise definition”

- **first used** at an informal meetup in **2009** in San Francisco  
(presentations from Voldemort, Cassandra, Dynomite, HBase, Hypertable, CouchDB, and MongoDB)

[Sadalage & Fowler: NoSQL Distilled, 2012]

# Just Another Temporary Trend?

- There have been **other trends** here before
  - **object** databases, XML databases, etc.
- **But NoSQL databases:**
  - are answer to **real practical problems** big companies have
  - are often developed by the **biggest players**
  - outside academia but based on **solid theoretical results**
    - e.g. old results on distributed processing
  - widely used

# The End of Relational Databases?

- **Relational databases are not going away**
  - are ideal for a lot of structured data, reliable, mature, etc.
- **RDBMS became one option for data storage**

**Polyglot persistence** – using different data stores in different circumstances [Sadalage & Fowler: NoSQL Distilled, 2012]

Two trends:

1. **NoSQL** databases **implement standard** RDBMS features
2. **RDBMS** are **adopting** NoSQL principles

# RDBMS vs NoSQL: At a glance

<b>RDBMS</b>	<b>NoSQL</b>
- Structured and organized data	- Stands for Not Only SQL
- Structured query language (SQL)	- No declarative query language
- Data and its relationships are stored in separate tables	- No predefined schema
- Data Manipulation Language, Data Definition Language	- Key-Value pair storage, Column Store, Document Store, Graph databases
- TightConsistency	- Eventual consistency rather ACID property
- ACID	- CAP Theorem

# CAP Theorem

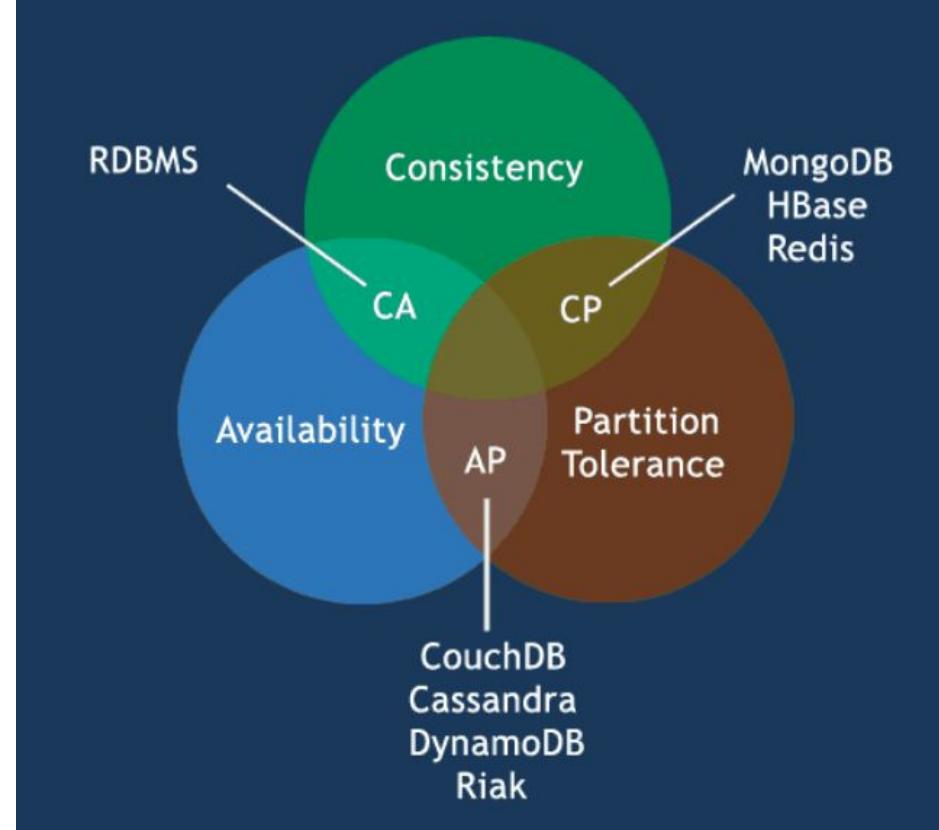
- **Consistency** : data in the database remains consistent after the execution of an operation e.g. after an update operation, all clients see the same data
  - Distributed System
- **Availability**: system is available, no downtime (emphasis in NoSQL)
- **Partition Tolerance**: system continues operation in case of failure (different nodes keep working)

# CAP Theorem

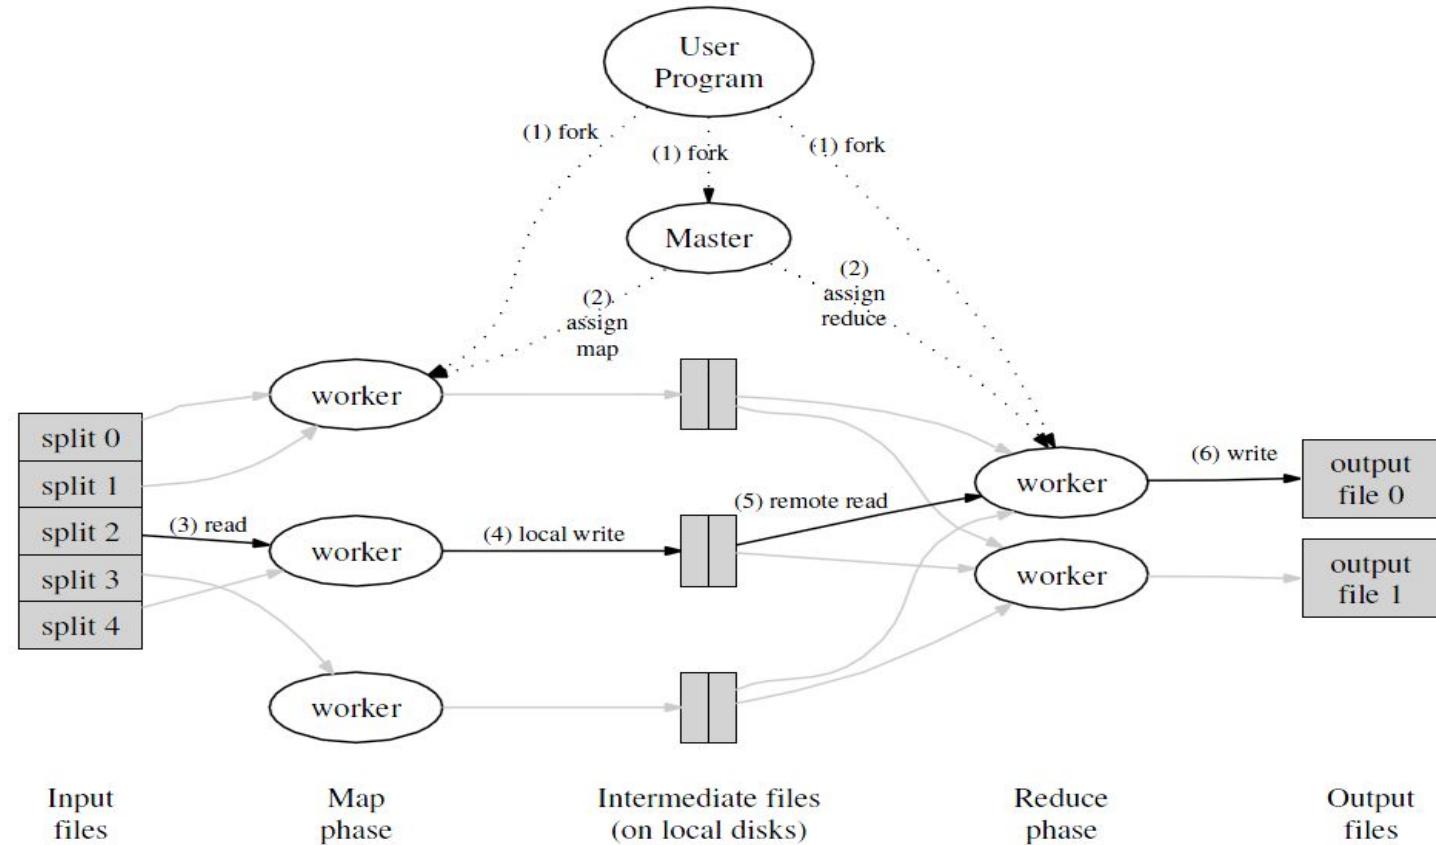
- **Consistency** : data in the database remains consistent after the execution of an operation e.g. after an update operation, all clients see the same data
  - Distributed System
- **Availability**: system is available, no downtime (emphasis in NoSQL)
- **Partition Tolerance**: system continues operation in case of failure (different nodes keep working)

Unlike ACID properties, it is very difficult to fulfill all three

# CAP Theorem



# Behind the scenes: Map Reduce



# Map Reduce: Features

- MapReduce is a **generic** approach for **distributed** processing of **large** data collections
- **Requires** a way to distribute the **data**
  - and to collect the results back after the processing
- The **user** must only specify two **functions**:  
**map & reduce**

# Map Reduce: Implementation

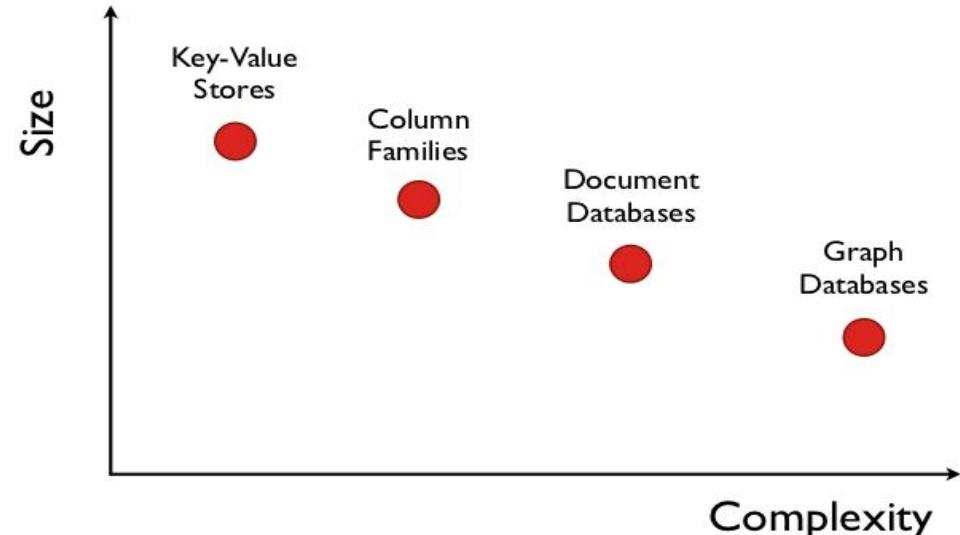


Amazon Elastic  
MapReduce



# Types of NoSQL Databases

- Key-value stores
- Column-family stores
- Document databases
- Graph databases



\* There is not a single solution that is better than the others

# Key-Value Stores

- A simple **hash table** (map), primarily used when all accesses to the database are via **primary key**
  - **key-value** mapping
- In RDBMS world: A table with two columns:
  - ID **column** (**primary key**)
  - DATA **column** storing the value (**unstructured BLOB**)
- **Basic operations:**
  - **Put** a value for a key                              `put(key, value)`
  - **Get** the value for the key                              `value := get(key)`
  - **Delete** a key-value                                      `delete(key)`

# Key-Value Stores: Architecture

1. Embedded systems
  - the system is a **library** and the DB runs **within your** system
2. Large-scale **Distributed** stores Architecture often as a **distributed hash table** (DHT)

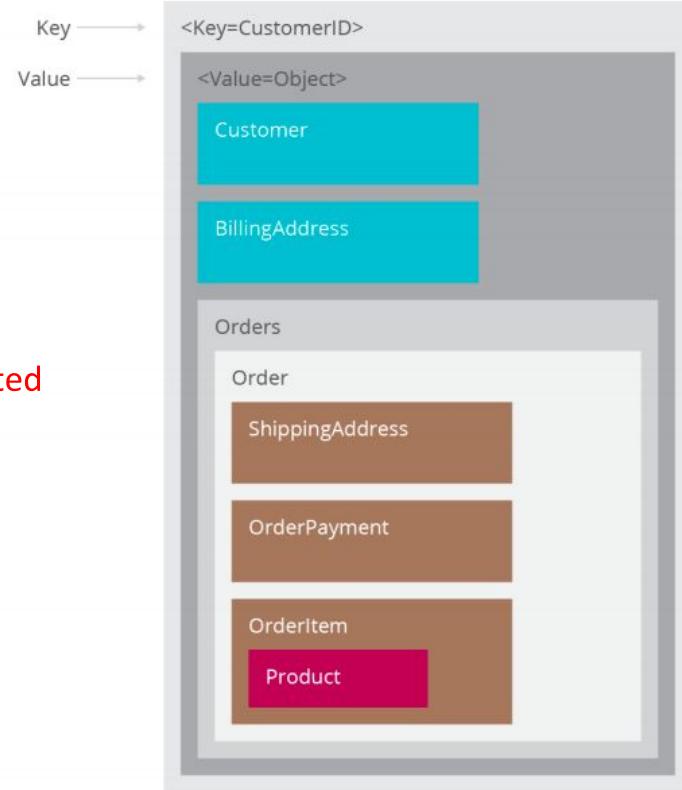
Features: it is **simple**

- great **performance**, easily scaled

# Key Value Stores

- A key may be strings, hashes, lists, sets, sorted sets
- Key-value stores can be used as collections, dictionaries, associative arrays etc.
- Key-value stores follow the **Availability** and **Partition Tolerance** aspects of CAP theorem
- Suitable for: shopping cart contents, individual values like color schemes

# Key Value Stores



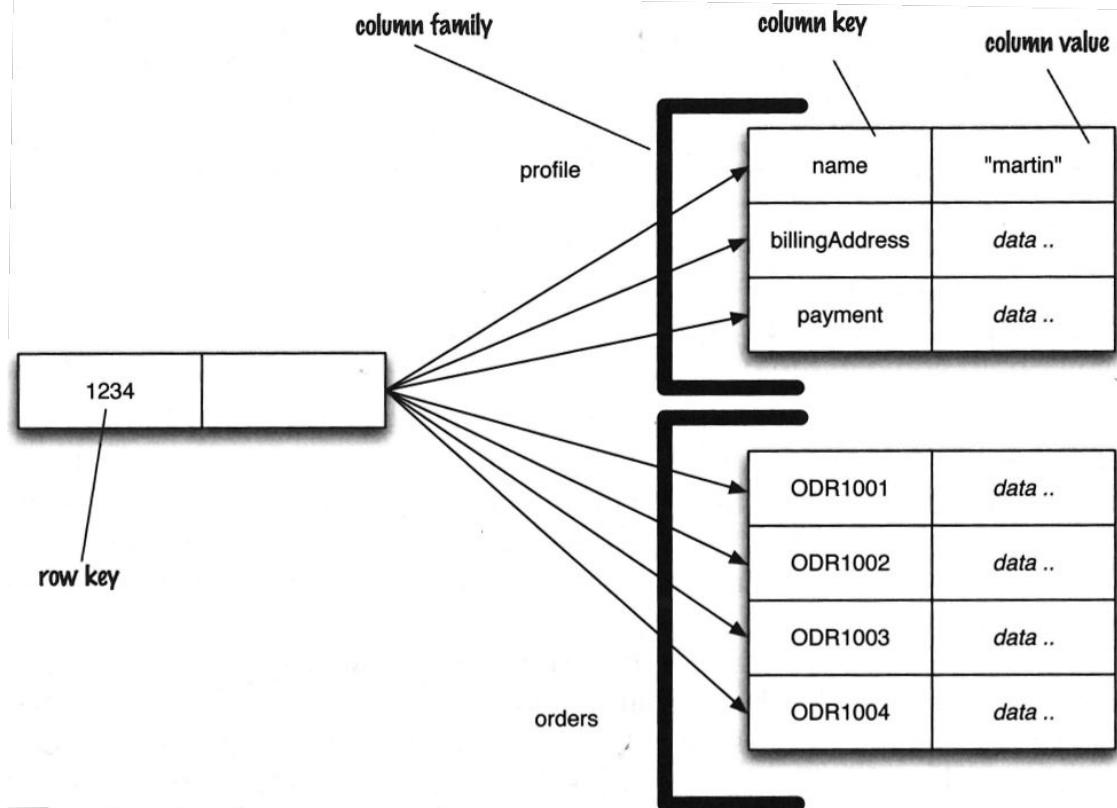
# Key-value Stores: Representatives



# Column-family Stores: Basics

- Data model: **rows** that have **many columns** associated with a **row key**
- **Column families** are groups of related data (**columns**) that are often **accessed together**
  - e.g., for a **customer** we typically access all **profile** information at the same time, but not customer's **orders**
- Best for \_\_\_\_\_ queries (MAX, MIN, COUNT, AVG etc.)

# Column-family Stores: Example



# Column-family Stores: Representatives



*Cassandra*

Google  
BigTable



HYPERTABLE

HBASE

ACCUMULO™

# Document Databases: Basics

- Basic concept of data: *Document*
- Documents are **self-describing** pieces of data
  - **Hierarchical tree** data **structures**
  - Nested associative arrays (maps), collections, scalars
  - XML, JSON (JavaScript Object Notation), BSON, ...
- Documents in a **collection** should be “similar”
  - Their **schema** can **differ**
- **Documents stored in the **value** part of key-value**
  - Key-value stores where the values are **examinable**
  - Building search **indexes** on various **keys/fields**

# Document Databases: Data Example

```
key=3 -> { "personID": 3,
    "firstname": "Martin",
    "likes": [ "Biking", "Photography" ],
    "lastcity": "Boston",
    "visited": [ "NYC", "Paris" ] }

key=5 -> { "personID": 5,
    "firstname": "Pramod",
    "citiesvisited": [ "Chicago", "London", "NYC" ],
    "addresses": [
        { "state": "AK",
            "city": "DILLINGHAM" },
        { "state": "MH",
            "city": "PUNE" } ],
    "lastcity": "Chicago" }
```

# Document Databases: Queries

## Example in MongoDB syntax

- **Query** language expressed via **JSON**
- clauses: where, sort, count, sum, etc.

SQL:           SELECT \* FROM users

MongoDB:   db.users.find()

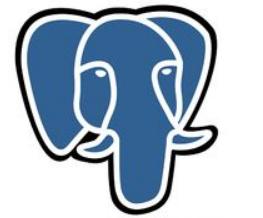
SELECT \* FROM users WHERE personID = 3  
db.users.find( { "personID": 3 } )

SELECT firstname, lastcity FROM users WHERE personID = 5  
db.users.find( { "personID": 5 }, {firstname:1, lastcity:1} )

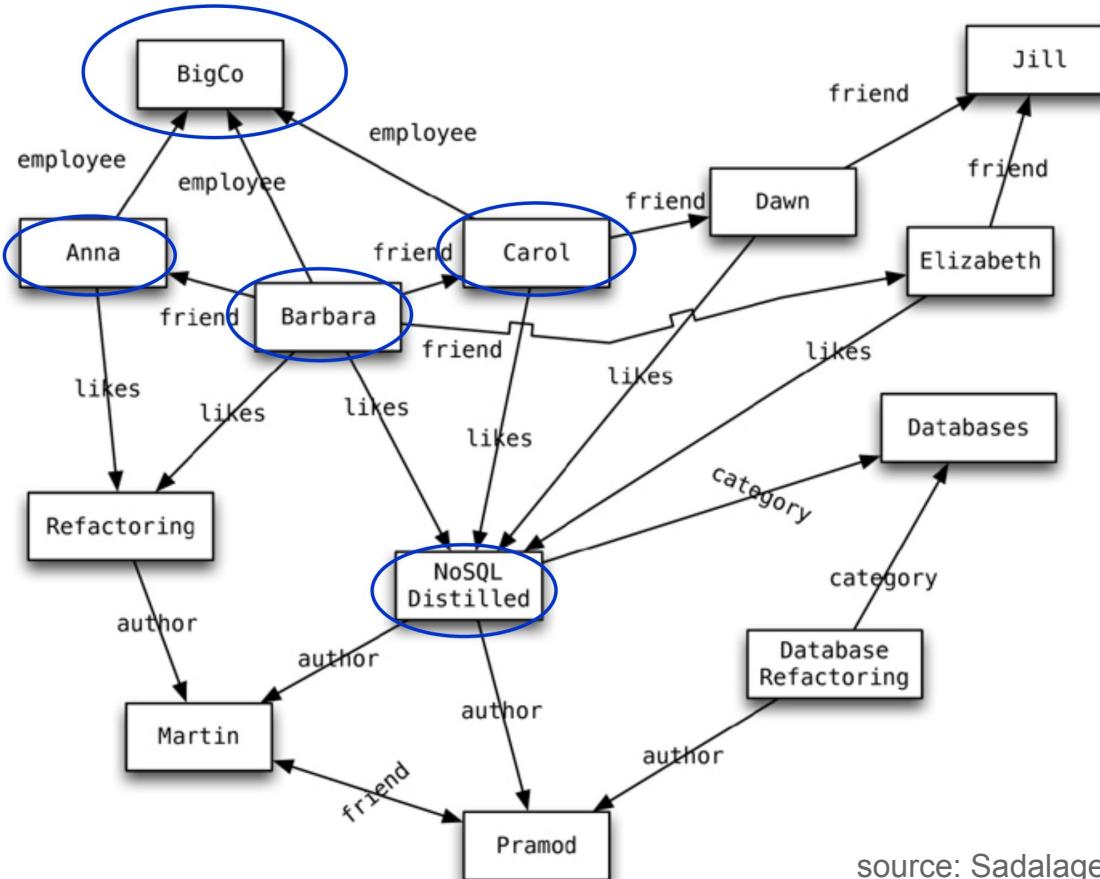
# Relational vs Document Model

Relational Model	Document Model
Tables	Collections
Rows	Documents
Columns	Key-value pairs
Joins	N/A

# Document Databases: Representatives



# Graph Databases: Example



Collection of nodes,  
edges and attributes

# Graph Databases: Representatives



# One Example: Facebook

## Facebook statistics (2016)

- **1.86 billion** monthly active users
- **4 million** 'likes' per minute
- **250 billion** stored photos (350 million uploaded daily)
- **300 PB** of user data stored (2014)

2009: 10,000 servers

2010: 30,000 servers

2012: 180,000 servers (estimated)



# Facebook: Database Tech. Behind

## Apache Hadoop <http://hadoop.apache.org/>



- **Hadoop File System (HDFS)**
  - over 100 PB in a single HDFS cluster
- an open source implementation of **MapReduce**:
  - Enables efficient calculations on massive amounts of data

## Apache Hive <http://hive.apache.org/>



- **SQL-like access** to Hadoop-stored data
- integration of **MapReduce** query evaluation

# Facebook: Database Tech. Behind (2)

## Apache HBase <http://hbase.apache.org/>

- a Hadoop **column-family** database
- used for e-mails, instant messaging and SMS
- **replacement** for MySQL and Cassandra



## Memcached <http://memcached.org/>

- distributed key-value store
- used as a **cache** between web servers and MySQL servers in the beginning of FB



# Facebook: Database Tech. Behind (3)

## Apache Giraph <http://giraph.apache.org/>

- **graph** database
- facebook **users and connections** is one very large graph
- used since 2013 for various analytic tasks (trillion edges)



## RocksDB <http://rocksdb.org/>

- high-performance **key-value store**
- developed **internally** in FB, now open-source





# Sharding

What is database sharding ?

# Sharding

- A database shard is a **horizontal partition** of data in a database or search engine.
- Each individual partition is referred to as a shard or database shard
- Each shard is held on a separate database server instance, to spread load.
- Very important for high availability

# Applications at C-DAC

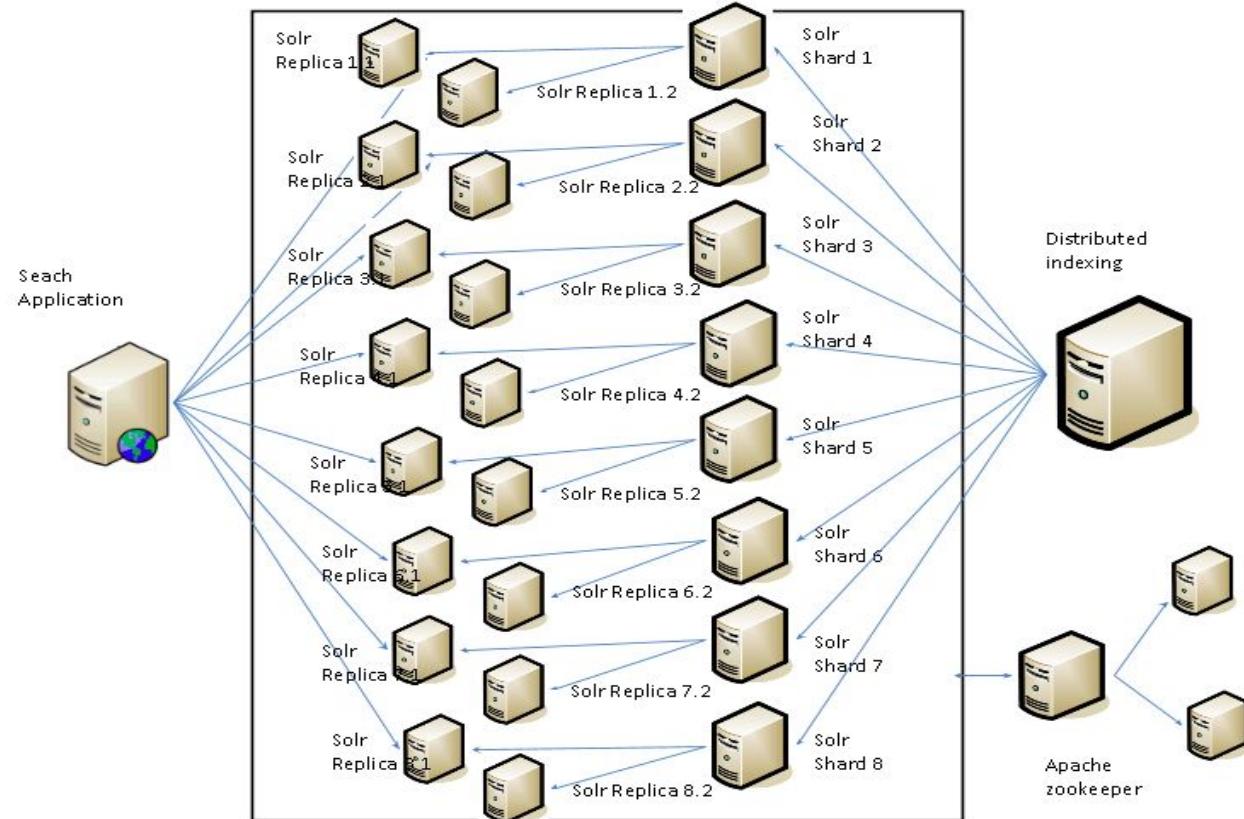
- National Electoral Search (election season !!!)
  - e-RaktKosh - A nationwide initiative to reach citizens
  - Hospital Management Information Systems (high availability report generation)

# National Electoral Search

- We have already discussed the objectives
  - High Availability
  - < 50ms response time
  - > 50 lakh concurrent users
  - Multiple search criteria
  - ...

# National Electoral Search

- 6 TB data + indexes + backup
- Envisaged infrastructure
  - 8 Servers
  - 64 GB RAM each



# National Electoral Search

How did we end up doing it ?

- 2 servers (8GB RAM each - commodity hardwares)
- Lucene (not publicized as NoSQL)
- The voter data was partitioned as per the states
- Indexes were created on each state's data
- For location mapping, a custom algorithm was developed

# Hospital Management System

Task: Generate patient reports for various investigations and stats

## Biochemistry

<b>CR No</b>	: 109111900354637	<b>Lab/Study No.</b>	: 0003	<b>Requisition Date</b> : 01-Apr-2019 10:17
<b>Patient Name</b>	: Eeeemmm	<b>Age/Sex</b>	: 34 Yr /M	<b>Coll./Study Date</b> : 01-Apr-2019 10:17
<b>Sample Type/No</b>	: Serum/0003	<b>Ward/OPD</b>	: OPD	<b>Reporting Date</b> : 01-Apr-2019 10:25
<b>Dept/Unit</b>	: General Medicine Casualty Unit			

<b>Investigation</b>	<b>Result</b>	<b>Unit</b>	<b>Ref. Range</b>
----------------------	---------------	-------------	-------------------

Kidney Function Test			
Serum Urea	<b>122</b>	mg/dL	13 - 43
Creatinine	<b>122</b>	mg/dl	0.7 - 1.3
Uric Acid	<b>122</b>	mg/dl	3.5 - 7.2
Colour			
Serum Calcium	<b>122</b>	mg/dL	8.6 - 10
Chloride	<b>87</b>	meq/l	98 - 107
File Upload	-		

**Comments:**

# Hospital Management System

## Challenges

- Data from 300 hospitals for 1000 patients daily
- Each patient has 500 records (average)
- 150 million records each day
- At peak time, 1 million records each second
- Highly unstructured - each patient has different tests and values

# Hospital Management System

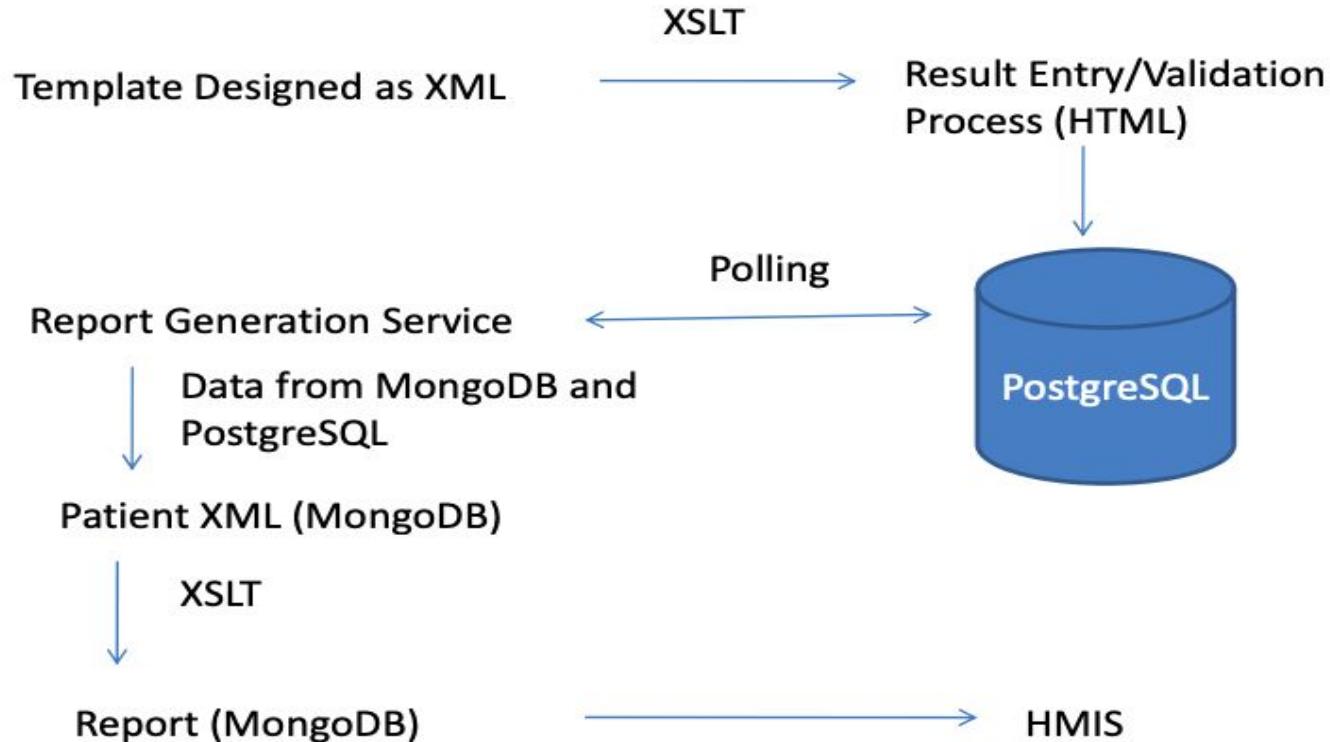
## Solution

- Utilize strength of both RDBMS and NoSQL
- RDBMS
  - Good at mapping test and values
- NoSQL
  - Good at storing and retrieving data

# Hospital Management System

- Utilize strength of both RDBMS and NoSQL
    - RDBMS
      - Good at mapping test and values
      - PostgreSQL
    - NoSQL
      - Good at storing and retrieving data
      - MongoDB GridFS
  - The concept is being used in upcoming Distributed Health System (DHS) for entire country

# Hospital Management System





सर्वीडैक  
CDCC

# e-RaktKosh ([www.eraktkosh.in](http://www.eraktkosh.in))

Older version ▾ Skip to main content ▾ A A+ A A- Screen Readers

GENERAL SERVICES DOWNLOADS ABOUT US LOGIN

## View Information

Search Your Location

Location

View Blood Banks

Login

My Donations

eRaktkosh Login

Register as Donor

### Locate Near By Blood Banks

Locate nearby blood banks or search for any location with our full featured map.

< >

Blood Cell - NHM Initiatives

96 Active Camps

1735 Active e-Blood Banks

30000+ Blood Units Available

Why Donate Blood

Get Mobile Apps

<https://eraktkosh.in/BLDAHIMS/bloodbank/transactions/bbpublicindex.html#myCarousel>

# Use Case

- Task: Find update frequency, who missed the update, outliers, discover hidden stats

# NoSQL Use Case

- 1800 Blood Banks across the country update blood stock daily
- Average update frequency is 4 times a day
- Each update inserts 120 records
  - 864000 records a day
  - 25 million records per month
- Need to be available to
  - Citizens (Portal, Mobile Apps, UMANG)
  - Ministers (CM Dashboards)
  - Blood banks (API)

# NoSQL Use Case

- Expected Approach
  - For each blood bank we create a **document**
  - The document contains
    - Details of Blood Bank
    - Blood Stock
    - Updation date and other meta data
  - Traverse through the document with aggregation queries

# NoSQL Use Case

- Expected Approach
  - For each blood bank we create a **document**
  - The document contains
    - Details of Blood Bank
    - Blood Stock
    - Updation date and other meta data
  - Traverse through the document with aggregation queries

Is it the optimal approach ?

# NoSQL Use Case

- Recall NoSQL schema design is application specific
- If statistics are required, better arrange them in such a form (high performance use case)

# NoSQL Use Case

- Approach followed
  - Construct another document where on each insert aggregate statistics
  - Can be done in RDBMS also, but data is distributed so RDBMS approach is not optimal



# Questions?

Please, any questions? Good question is a gift...

*MongoDB (from “humongous”) is a scalable,  
high-performance, open source, schema-free,  
document-oriented database.*

-- Mongodb.org

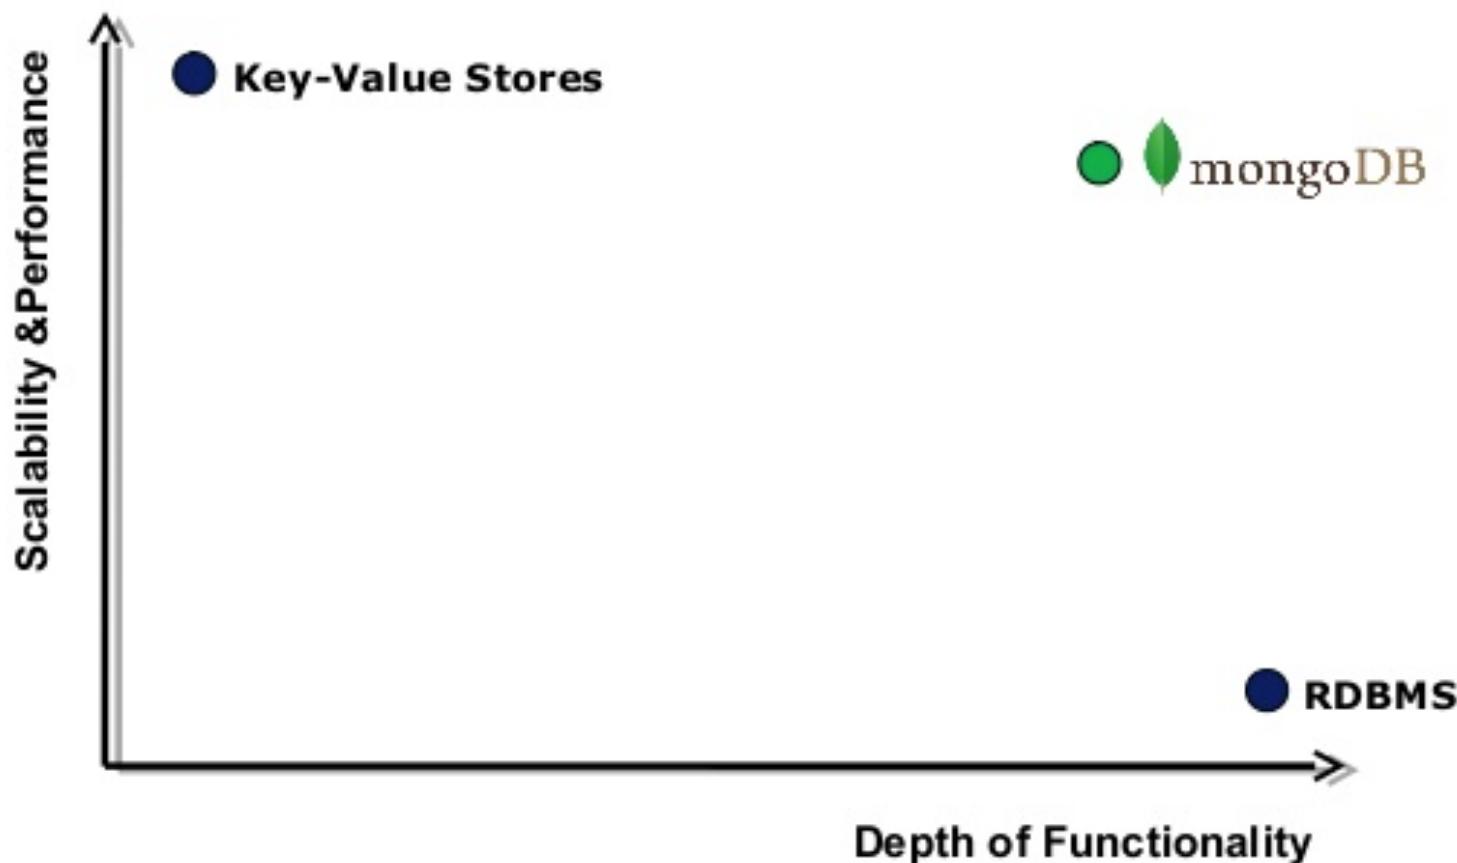


mongoDB

# Background

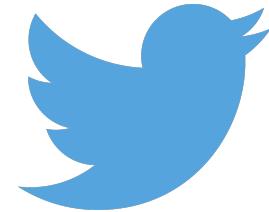
- A NoSQL database of type document oriented
- Eschews the traditional table-based relational database structure in favor of JSON-like documents with dynamic schemas (MongoDB calls the format BSON)
- First developed by MongoDB Inc in October 2007
- Shifted to open source in 2009

# Where MongoDB Stands?



# Examples of JSON Format

```
"filter_level":"medium",
"contributors":null,
"text":"Do you think neymar will score his first goal tonight ???",
"geo":{
    "type":"Point",
    "coordinates":[30.000000,20.000000]
},
"retweeted":false,
"in_reply_to_screen_name":null,
"truncated":false,
"lang":"en",
"entities":{
    "symbols":[],
    "urls":[],
    "hashtags":[],
    "user_mentions":[]
},
"in_reply_to_status_id_str":null,
"id":363356918,
"source":<a href=\"http://twitter.com/download/android\" rel=\"nofollow\">Twitter for Android</a>",
"in_reply_to_user_id_str":null,
"favorited":false,
"in_reply_to_status_id":null,
"retweet_count":0,
"created_at":"Fri Aug 02 17:53:34 +0000 2013",
"in_reply_to_user_id":null,
"favorite_count":0,
"id_str":"363356918",
"place":null,
"user":{
    "location":"",
    "default_profile":false,
    "profile_background_tile":true,
    "statuses_count":8100,
    "lang":"en",
    "profile_link_color":"009989",
    "profile_banner_url":"https://pbs.twimg.com/profile_banners/414336796/1379000000",
    "id":414336796,
    "following":null,
    "protected":false,
    "favourites_count":855,
    "profile_text_color":"3c3940",
    "description":"My dream is all my life"
}
```



# MongoDB Basics (cont.)



Each document within a collection can have its own unique set of fields

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value  
← field: value  
← field: value  
← field: value

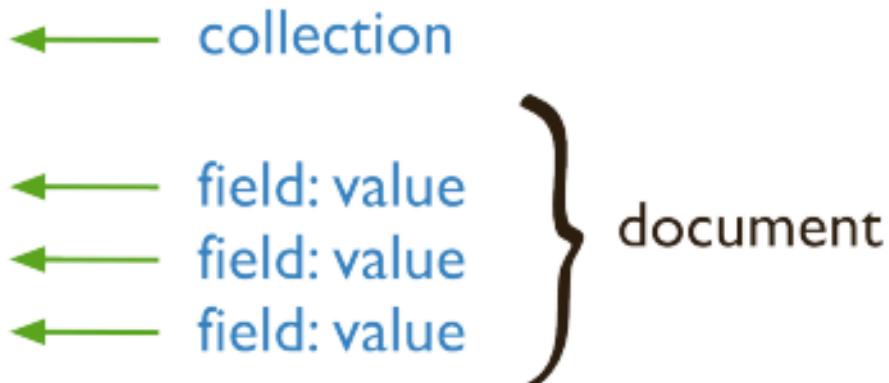
# MongoDB CRUD

- MongoDB provides rich semantics for reading and manipulating data
- CRUD = Create, Read, Update, and Delete

# CRUD: Create

In MongoDB

```
db.users.insert ( ← collection
  {
    name: "sue", ← field: value
    age: 26, ← field: value
    status: "A" ← field: value
  }
)
```



In SQL

```
INSERT INTO users ← table
          ( name, age, status ) ← columns
VALUES      ( "sue", 26, "A" ) ← values/row
```

# CRUD: Read

In MongoDB

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
).limit(5)
```

← collection  
← query criteria  
← projection  
← cursor modifier

In SQL

```
SELECT _id, name, address  
FROM   users  
WHERE  age > 18  
LIMIT  5
```

← projection  
← table  
← select criteria  
← cursor modifier

# CRUD: Update

In MongoDB

```
db.users.update(  
    { age: { $gt: 18 } },           ← collection  
    { $set: { status: "A" } },      ← update criteria  
    { multi: true }               ← update action  
)  
                                ← update option
```

In SQL

```
UPDATE users                ← table  
SET status = 'A'             ← update action  
WHERE age > 18              ← update criteria
```

# CRUD: Delete

In MongoDB

```
db.users.remove(  
    { status: "D" }  
)
```



← collection  
← remove criteria

In SQL

```
DELETE FROM users  
WHERE status = 'D'
```



← table  
← delete criteria

# Using MongoDB

- You can either install MongoDB on your machine, or visit:

[http://www.tutorialspoint.com/  
mongodb\\_terminal\\_online.php](http://www.tutorialspoint.com/mongodb_terminal_online.php)

# To Get Started...

- Global commands: help, exit, etc.
- Commands execute against the current database are executed against the db object, for example:
  - db.help(): returns a list of commands that you can use against db object
  - Note: db.help without () gives you the method body

# Create Database

- To create a wonderland database:

```
use wonderland
```

\* creates the database and switches to it

- To get the collections in the current database:

```
db.getCollectionNames()
```

# Insert Data

- To insert a document into the collection:

```
db.unicorns.insert(  
    {name: 'Aurora',  
     gender: 'f',  
     weight: 450}  
)
```

- \* Try out db.getCollectionNames() now, you'll see:

# List Documents in a Collection

- Try out:

```
db.unicorns.find()
```

- One more field is added: \_id
  - Every document must have a unique \_id field
  - Can generate your own or have MongoDB generate automatically for you

Remove all data: db.unicorns.remove({}).  
Get the data from: <http://bit.ly/iiitsdbms>

# Query Selector (cont.)

- Use  
 $\{field: value\}$   
to select documents that match the condition.
- If matching multiple conditions is desired:  
 $\{field1: value1, field2: value2...\}$   
\* This implies the \_\_\_\_\_ statement

# Comparison Operators in MongoDB

- `$lt` less than
- `$lte` less than or equal to
- `$gt` greater than
- `$gte` greater than or equal to
- `$ne` not equal to

(Q1) Find the male unicorns weigh more than 700 pounds

Ans1:

```
db.unicorns.find({gender: 'm', weight:  
{$gt: 700}})
```

(Q2) Find the unicorns that have no vampire field

**Ans 2:**

```
db.unicorns.find({ vampires: {$exists: false}})
```

(Q3) Find the unicorns that like apples or oranges

Ans 3:

```
db.unicorns.find({ loves: {$in: ['apple','orange']} })
```

(Q4) Find the female unicorns that either love apples or weigh less than 500 pounds

Ans 4:

```
db.unicorns.find({gender: 'f', $or: [{loves: 'apple'}, {weight: {$lt: 500}}]})
```

# CRUD: Update

- Intuitively, updating unicorn Roooooodles' weight to 590 can be:

```
db.unicorns.update(  
  {name: "Roooooodles"},  
  {weight: 590})
```

- But if you try:

```
db.unicorns.find({name: "Roooooodles"})
```

the result will be: \_\_\_\_\_

# CRUD: Update (cont.)

- The reason that no document was found was because the second parameter we supplied didn't have any update operators
- Therefore, the original document was replaced
- Try the following command to see:

```
db.unicorns.find({weight: 590})
```

# CRUD: Update (cont.)

- To fix the problem, we should do:

```
db.unicorns.update({weight: 590},  
{$set: {name: "Roooooodles",  
        dob: new Date(1979, 7, 18, 18, 44),  
        loves: ["apple"],  
        gender: "m",  
        vampires: 99}})
```

# CRUD: Update (cont.)

- The correct way to update at the beginning should therefore be:

```
db.unicorns.update({name: "Roooooodles"},  
{$set: {weight: 590}})
```

# More Update Operators

- **\$inc**: increment a field by a certain positive or negative amount
- **\$push**: add a value to the existing field

(Q5) Decrease unicorn Pilot's number of vampires by 2

Ans 5:

```
db.unicorns.update({name: 'Pilot'}, {$inc: {vampires: -2}})
```

(Q6) Add “sugar” to the list of food  
unicorn Aurora loves to eat

Ans 6:

```
db.unicorns.update({name: 'Aurora'}, {$push: {loves: 'sugar'}})
```

# Projection

- `find()` can take a second argument, which is the **project list**

- Example:

```
db.unicorns.find({}, {name:1, _id:0})
```

- The values following field names are boolean:
  - 1 means including the field
  - 0 means excluding the field
- Note that except excluding `_id`, the list cannot have a mixture of exclusion and inclusion

# Upserts

- An upsert updates the document if found or inserts it if not
- To enable upserting we pass a third parameter to update {upsert: true}

# Upserts (cont.)

- This will not do anything:

```
db.unicorns.update({name: "Walala"},  
{$inc: {vampires: 1}})
```

- Instead, do this:

```
db.unicorns.update({name: "Walala"},  
{$inc: {vampires: 1}},  
{upsert: true})
```

# Multiple Updates

- By default, update will only update a single document. Passing the third parameter {multi: true} will enable the multiple update

(Q7) Give all of the unicorns vaccine  
(set vaccinated to be true)

Ans 7:

```
db.unicorns.update({}, {$set: {vaccinated: true }},  
{multi:true});
```

```
db.unicorns.find({vaccinated: true});
```

(Q8) Sort the unicorns based on weights decreasingly

Ans 8:

```
db.unicorns.find().sort({weight: -1})
```

(Q9) Sort the unicorns based on the names increasingly, then the number of vampires decreasingly

Ans 9:

```
db.unicorns.find().sort({name: 1, vampires: -1})
```

(Q10) Get the second and third  
heaviest unicorns

Ans: 10

```
db.unicorns.find() .sort({weight: -1}) .limit(2) .skip(1)
```

(Q11) Count the number of unicorns  
who have more than 50 vampires

Ans 11:

```
db.unicorns.count({vampires: {$gt: 50}})
```

# References

- Karl Seguin, *The Little MongoDB Book*,  
<http://openmymind.net/mongodb.pdf>
- Kristina Chodorow, *MongoDB: The Definite Guide*, O'Reilly
- MongoDB CRUD Operations,  
<https://docs.mongodb.org/master/MongoDB-crud-guide-master.pdf>

# MongoDb(No SQL Database)

## Questions

Structure of file “[restaurants.json](#)”

```
{  
  "address": {  
    "building": "1007",  
    "coord": [ -73.856077, 40.848447 ],  
    "street": "Morris Park Ave",  
    "zipcode": "10462"  
  },  
  "borough": "Bronx",  
  "cuisine": "Bakery",  
  "grades": [  
    { "date": { "$date": 1393804800000 }, "grade": "A", "score": 2 },  
    { "date": { "$date": 1378857600000 }, "grade": "A", "score": 6 },  
    { "date": { "$date": 1358985600000 }, "grade": "A", "score": 10 },  
    { "date": { "$date": 1322006400000 }, "grade": "A", "score": 9 },  
    { "date": { "$date": 1299715200000 }, "grade": "B", "score": 14 }  
  ],  
  "name": "Morris Park Bake Shop",  
  "restaurant_id": "30075445"  
}
```

- 1. Write a MongoDB query to display all the documents in the collection restaurants**
- 2. Write a MongoDB query to display the fields restaurant\_id, name, borough and cuisine for all the documents in the collection restaurant.**
- 3. Write a MongoDB query to display the fields restaurant\_id, name, borough and cuisine, but exclude the field \_id for all the documents in the collection restaurant.**

- 4. Write a MongoDB query to display the fields restaurant\_id, name, borough and zip code, but exclude the field \_id for all the documents in the collection restaurant.**
- 5. Write a MongoDB query to display all the restaurant which is in the borough Bronx.**
- 6. Write a MongoDB query to display the first 5 restaurant which is in the borough Bronx.**
- 7. Write a MongoDB query to display the next 5 restaurants after skipping first 5 which are in the borough Bronx.**
- 8. Write a MongoDB query to find the restaurants who achieved a score more than 90.**
- 9. Write a MongoDB query to find the restaurants that achieved a score, more than 80 but less than 100.**
- 10. Write a MongoDB query to find the restaurants which locate in latitude value less than -95.754168.**
- 11. Write a MongoDB query to find the restaurants that do not prepare any cuisine of 'American' and their grade score more than 70 and latitude less than -65.754168.**
- 12. Write a MongoDB query to find the restaurants which do not prepare any cuisine of 'American' and achieved a score more than 70 and located in the longitude less than -65.754168.**

**Note : Do this query without using \$and operator.**

- 13. Write a MongoDB query to find the restaurants which do not prepare any cuisine of 'American' and achieved a grade point 'A' not belongs to the borough Brooklyn. The document must be displayed according to the cuisine in descending order.**
- 14. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which contain 'Wil' as first three letters for its name.**
- 15. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which contain 'ces' as last three letters for its name.**
- 16. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which contain 'Reg' as three letters somewhere in its name.**
- 17. Write a MongoDB query to find the restaurants which belong to the borough Bronx and prepared either American or Chinese dish.**
- 18. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which belong to the borough Staten Island or Queens or Bronxor Brooklyn.**
- 19. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which are not belonging to the borough Staten Island or Queens or Bronxor Brooklyn.**
- 20. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which achieved a score which is not more than 10.**

- 21. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which prepared dish except 'American' and 'Chinees' or restaurant's name begins with letter 'Wil'.**
- 22. Write a MongoDB query to find the restaurant Id, name, and grades for those restaurants which achieved a grade of "A" and scored 11 on an ISODate "2014-08-11T00:00:00Z" among many of survey dates.**
- 23. Write a MongoDB query to find the restaurant Id, name and grades for those restaurants where the 2nd element of grades array contains a grade of "A" and score 9 on an ISODate "2014-08-11T00:00:00Z".**
- 24. Write a MongoDB query to find the restaurant Id, name, address and geographical location for those restaurants where 2nd element of coord array contains a value which is more than 42 and upto 52.**
- 25. Write a MongoDB query to arrange the name of the restaurants in ascending order along with all the columns.**
- 26. Write a MongoDB query to arrange the name of the restaurants in descending order along with all the columns.**
- 27. Write a MongoDB query to arranged the name of the cuisine in ascending order and for that same cuisine borough should be in descending order.**
- 28. Write a MongoDB query to know whether all the addresses contains the street or not.**
- 29. Write a MongoDB query which will select all documents in the restaurants collection where the coord field value is Double.**

**30. Write a MongoDB query which will select the restaurant Id, name and grades for those restaurants which returns 0 as a remainder after dividing the score by 7.**

**31. Write a MongoDB query to find the restaurant name, borough, longitude and attitude and cuisine for those restaurants which contains 'mon' as three letters somewhere in its name.**

**32. Write a MongoDB query to find the restaurant name, borough, longitude and latitude and cuisine for those restaurants which contain 'Mad' as first three letters of its name.**